

## Data Cleaning

```
import pandas as pd
import numpy as np

# Sample dataset with missing values and duplicates
data1 = pd.DataFrame({
    'ID': [1, 2, 3, 4, 4],
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'David'],
    'Age': [25, 30, np.nan, 40, 40],
    'City': ['New York', 'Los Angeles', 'Chicago', 'Houston', 'Houston']
})

print("Original Data:")
print(data1)

# Remove duplicate rows
data1 = data1.drop_duplicates()

# Fill missing values in 'Age' with the column mean
data1['Age'].fillna(data1['Age'].mean(), inplace=True)

print("\nCleaned Data:")
print(data1)
```

Original Data:

	ID	Name	Age	City
0	1	Alice	25.0	New York
1	2	Bob	30.0	Los Angeles
2	3	Charlie	NaN	Chicago
3	4	David	40.0	Houston
4	4	David	40.0	Houston

Cleaned Data:

	ID	Name	Age	City
0	1	Alice	25.000000	New York
1	2	Bob	30.000000	Los Angeles
2	3	Charlie	31.666667	Chicago
3	4	David	40.000000	Houston

<ipython-input-1-8f580b932d8f>:19: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained ass The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
data1['Age'].fillna(data1['Age'].mean(), inplace=True)
```

## Data Integration

```
# Second dataset to merge
data2 = pd.DataFrame({
    'ID': [1, 2, 3, 5],
    'Income': [50000, 60000, 55000, 45000],
    'Gender': ['F', 'M', 'M', 'F']
})

# Merge datasets on 'ID'
merged_data = pd.merge(data1, data2, on='ID', how='left')

print("Merged Data:")
print(merged_data)
```

Merged Data:

	ID	Name	Age	City	Income	Gender
0	1	Alice	25.000000	New York	50000.0	F
1	2	Bob	30.000000	Los Angeles	60000.0	M
2	3	Charlie	31.666667	Chicago	55000.0	M
3	4	David	40.000000	Houston	NaN	NaN

## Data Transformation

```
# Encode categorical variable 'Gender' using one-hot encoding
merged_data = pd.get_dummies(merged_data, columns=['Gender'], drop_first=True)

# Normalize 'Income' column using Min-Max scaling
merged_data['Income_Normalized'] = (
```

```
(merged_data['Income'] - merged_data['Income'].min()) /
(merged_data['Income'].max() - merged_data['Income'].min())
)

print("Transformed Data:")
print(merged_data)
```

Transformed Data:

	ID	Name	Age	City	Income	Gender_M	Income_Normalized
0	1	Alice	25.000000	New York	50000.0	False	0.0
1	2	Bob	30.000000	Los Angeles	60000.0	True	1.0
2	3	Charlie	31.666667	Chicago	55000.0	True	0.5
3	4	David	40.000000	Houston	NaN	False	NaN

## Data Reduction

```
# Drop columns that are not useful for analysis
reduced_data = merged_data.drop(columns=['City', 'Income'])

print("Reduced Final Data:")
print(reduced_data)
```

Reduced Final Data:

	ID	Name	Age	Gender_M	Income_Normalized
0	1	Alice	25.000000	False	0.0
1	2	Bob	30.000000	True	1.0
2	3	Charlie	31.666667	True	0.5
3	4	David	40.000000	False	NaN

## Upload the Excel File to Google Colab

```
from google.colab import files
uploaded = files.upload()
```

Choose Files student\_data.xlsx

- **student\_data.xlsx**(application/vnd.openxmlformats-officedocument.spreadsheetml.sheet) - 13912 bytes, last modified: 4/6/2025 - 100% done

Saving student\_data.xlsx to student\_data.xlsx

```
import pandas as pd
```

```
# Load the Excel file into a DataFrame
df = pd.read_excel("student_data.xlsx")
```

```
# Display the data
print(df)
```

StudentID Name Age Gender Grade Email \

0	1000	Jerry	NaN	Other	C	<a href="mailto:zhill@gmail.com">zhill@gmail.com</a>
1	1001	Erica	NaN	NaN	D	<a href="mailto:yayala@walker.com">yayala@walker.com</a>
2	1002	Ashley	NaN	F	C	<a href="mailto:brian89@gmail.com">brian89@gmail.com</a>
3	1003	Michael	NaN	Other	A	<a href="mailto:hickmanrebecca@yahoo.com">hickmanrebecca@yahoo.com</a>
4	1004	Eileen	20.0	Other	B	<a href="mailto:michaelbaker@fernandez-davis.net">michaelbaker@fernandez-davis.net</a>
..	...	...	...	...	...	...
100	1022	Danielle	18.0	F	NaN	<a href="mailto:delacruzbarry@montes.net">delacruzbarry@montes.net</a>
101	1080	Jack	NaN	NaN	NaN	<a href="mailto:ywatkins@kirk-peters.com">ywatkins@kirk-peters.com</a>
102	1053	Amanda	24.0	F	D	<a href="mailto:juliebailey@yahoo.com">juliebailey@yahoo.com</a>
103	1061	Justin	NaN	F	F	<a href="mailto:bryanking@quinn.com">bryanking@quinn.com</a>
104	1072	Gregory	NaN	F	F	<a href="mailto:kimberlymcconnell@hotmail.com">kimberlymcconnell@hotmail.com</a>

	City	EnrollmentDate
0	Los Angeles	NaT
1	Los Angeles	2024-03-15
2	Los Angeles	2025-03-20
3	Phoenix	2024-01-28
4	New York	2024-07-12
..	...	...
100	NaN	NaT
101	NaN	2023-11-26
102	Los Angeles	NaT
103	Chicago	NaT
104	Los Angeles	NaT

```
[105 rows x 8 columns]
```

```
df.head()
```

	StudentID	Name	Age	Gender	Grade	Email	City	EnrollmentDate	
0	1000	Jerry	NaN	Other	C	zhill@gmail.com	Los Angeles	NaT	
1	1001	Erica	NaN	NaN	D	yayala@walker.com	Los Angeles	2024-03-15	
2	1002	Ashley	NaN	F	C	brian89@gmail.com	Los Angeles	2025-03-20	
3	1003	Michael	NaN	Other	A	hickmanrebecca@yahoo.com	Phoenix	2024-01-28	
4	1004	Eileen	20.0	Other	B	michaelbaker@fernandez-davis.net	New York	2024-07-12	

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

Data Cleaning

```
# Make a copy of the dataset for cleaning
cleaned_df = df.copy()

# Remove duplicate rows
cleaned_df = cleaned_df.drop_duplicates()

# Fill missing values
cleaned_df['Age'] = cleaned_df['Age'].fillna(cleaned_df['Age'].mean())
cleaned_df['Gender'] = cleaned_df['Gender'].fillna(cleaned_df['Gender'].mode()[0])
cleaned_df['Grade'] = cleaned_df['Grade'].fillna(cleaned_df['Grade'].mode()[0])
cleaned_df['City'] = cleaned_df['City'].fillna(cleaned_df['City'].mode()[0])
cleaned_df['EnrollmentDate'] = cleaned_df['EnrollmentDate'].fillna(pd.Timestamp('2023-01-01'))

# View cleaned data
cleaned_df.head()
```

	StudentID	Name	Age	Gender	Grade	Email	City	EnrollmentDate	
0	1000	Jerry	21.860465	Other	C	zhill@gmail.com	Los Angeles	2023-01-01	
1	1001	Erica	21.860465	F	D	yayala@walker.com	Los Angeles	2024-03-15	
2	1002	Ashley	21.860465	F	C	brian89@gmail.com	Los Angeles	2025-03-20	
3	1003	Michael	21.860465	Other	A	hickmanrebecca@yahoo.com	Phoenix	2024-01-28	
4	1004	Eileen	20.000000	Other	B	michaelbaker@fernandez-davis.net	New York	2024-07-12	

Next steps: [Generate code with cleaned\\_df](#) [View recommended plots](#) [New interactive sheet](#)

Data Integration

```
# Create simulated additional data for merging
import random

additional_data = pd.DataFrame({
    'StudentID': cleaned_df['StudentID'].sample(frac=0.8).values,
    'Club': np.random.choice(['Science', 'Art', 'Sports', 'None'], size=int(0.8 * len(cleaned_df)))
})

# Merge on StudentID
merged_df = pd.merge(cleaned_df, additional_data, on='StudentID', how='left')

# View merged data
merged_df.head()
```

	StudentID	Name	Age	Gender	Grade	Email	City	EnrollmentDate	Club	
0	1000	Jerry	21.860465	Other	C	zhill@gmail.com	Los Angeles	2023-01-01	None	
1	1001	Erica	21.860465	F	D	yayala@walker.com	Los Angeles	2024-03-15	NaN	
2	1002	Ashley	21.860465	F	C	brian89@gmail.com	Los Angeles	2025-03-20	NaN	
3	1003	Michael	21.860465	Other	A	hickmanrebecca@yahoo.com	Phoenix	2024-01-28	NaN	
4	1004	Eileen	20.000000	Other	B	michaelbaker@fernandez-davis.net	New York	2024-07-12	Art	

Next steps: [Generate code with merged\\_df](#) [View recommended plots](#) [New interactive sheet](#)

Data Transformation


```
# Make a copy for transformation
transformed_df = merged_df.copy()

# One-hot encode categorical columns
transformed_df = pd.get_dummies(transformed_df, columns=['Gender', 'Grade', 'Club'], drop_first=True)

# Normalize Age (Min-Max Scaling)
transformed_df['Age_Normalized'] = (
    (transformed_df['Age'] - transformed_df['Age'].min()) /
    (transformed_df['Age'].max() - transformed_df['Age'].min())
)

# Convert EnrollmentDate to datetime
transformed_df['EnrollmentDate'] = pd.to_datetime(transformed_df['EnrollmentDate'])

# View transformed data
transformed_df.head()
```




	StudentID	Name	Age	Email	City	EnrollmentDate	Gender_M	Gender_Other	Grade_B	Grade_C	Grade_D
0	1000	Jerry	21.860465	zhill@gmail.com	Los Angeles	2023-01-01	False	True	False	True	False
1	1001	Erica	21.860465	yayala@walker.com	Los Angeles	2024-03-15	False	False	False	False	True
2	1002	Ashley	21.860465	brian89@gmail.com	Los Angeles	2025-03-20	False	False	False	True	False
3	1003	Michael	21.860465	hickmanrebecca@yahoo.com	Phoenix	2024-01-28	False	True	False	False	False
4	1004	Eileen	20.000000	michaelbaker@fernandez-davis.net	New York	2024-07-12	False	True	True	False	False

Next steps: [Generate code with transformed\\_df](#) [View recommended plots](#) [New interactive sheet](#)

Data Reduction

```
# Drop irrelevant or less useful columns
reduced_df = transformed_df.drop(columns=['Email', 'City', 'EnrollmentDate'])

# Final preprocessed dataset
reduced_df.head()
```



	StudentID	Name	Age	Gender_M	Gender_Other	Grade_B	Grade_C	Grade_D	Grade_F	Club_None	Club_Science	Club_Sports	Ag
0	1000	Jerry	21.860465	False	True	False	True	False	False	True	False	False	
1	1001	Erica	21.860465	False	False	False	False	True	False	False	False	False	
2	1002	Ashley	21.860465	False	False	False	True	False	False	False	False	False	
3	1003	Michael	21.860465	False	True	False	False	False	False	False	False	False	
4	1004	Eileen	20.000000	False	True	True	False	False	False	False	False	False	

Next steps: [Generate code with reduced\\_df](#) [View recommended plots](#) [New interactive sheet](#)